

# RAGSuite | PRD v1.0

**Document Version:** 1.0

**Last Updated:** September 23, 2025

**Status:** Draft for Review

---

## 0) Scope Clarification — All-in-One App

- The platform is **one application** exposing both **Search** and **Chat** via the **same RAG pipeline**.
  - Integrators can enable **search**, **chat**, or **both** per embed/config (`mode: "search" | "chat" | "auto"`).
- 

## 1) Executive Summary

### 1.1 Product Overview

The AI Search & Chatbot Platform is an API-first, self-hosted solution that enables organizations to deploy intelligent search and conversational AI across their digital properties. The platform combines automated website crawling, retrieval-augmented generation (RAG), and local large language models to provide contextually accurate responses based on organizational content.

### 1.2 Business Objectives


- **Primary Goal:** Deliver a production-ready AI search and chatbot solution for on-premises deployment
- **Market Position:** Cost-effective alternative to cloud-based AI solutions with complete data control
- **Target Users:** Organizations requiring private AI deployments with custom content integration

### 1.3 Success Metrics


- **Technical:** <10 second response times, 99% uptime during business hours
  - **User Adoption:** 80%+ query satisfaction rate within first month of deployment
  - **Business:** Complete v1 delivery within 16-week development cycle
-

## 2) Product Architecture


### 2.1 System Architecture (Context)



### 2.2 Container/Deployment Topology



## 2.3 Backend Components (C4-L3)



## 2.4 Technology Stack

### Backend API Services

- **Framework:** FastAPI (Python 3.9+)
- **Database:** PostgreSQL 14+
- **Vector Database:** Chroma
- **Cache:** Redis
- **Web Crawler:** Scrapy with Playwright headless rendering
- **LLM Integration:** Ollama with Mistral 7B (default) or Llama 2 7B
- **Embeddings:** Sentence Transformers (all-MiniLM-L6-v2, 384-d)

### Frontend Applications

- **Framework:** React 18+ with TypeScript
- **State Management:** Redux Toolkit or Zustand
- **UI Library:** Material-UI or Tailwind CSS
- **Build Tool:** Vite
- **API Client:** Axios with OpenAPI code generation

### Infrastructure

- **Containerization:** Docker & Docker Compose
- **Process Management:** systemd or PM2 (optional)
- **Web Server:** Nginx (reverse proxy)
- **Monitoring:** Basic logging with structured JSON

## 3) API Specification

### 3.1 Core API Endpoints

#### Authentication & Configuration

- POST /api/v1/auth/login - Admin authentication
- GET /api/v1/config/system - System configuration
- PUT /api/v1/config/system - Update system settings
- GET /api/v1/config/themes - Available themes
- PUT /api/v1/config/branding - Update branding settings
- GET /api/v1/config/rag-defaults - Read RAG/reranker defaults

#### Content Management

- POST /api/v1/crawl/sites - Add crawling target
- GET /api/v1/crawl/sites - List crawling targets
- PUT /api/v1/crawl/sites/{id} - Update crawl configuration
- DELETE /api/v1/crawl/sites/{id} - Remove crawling target
- POST /api/v1/crawl/start/{id} - Start crawling job
- GET /api/v1/crawl/status/{id} - Check crawling status
- PUT /api/v1/crawl/preview - Return render+text sample for a URL (headless)

#### Document Management

- GET /api/v1/documents - List indexed documents
- POST /api/v1/documents/upload - Upload document manually
- PUT /api/v1/documents/{id} - Update document metadata
- DELETE /api/v1/documents/{id} - Remove document from index
- GET /api/v1/documents/{id}/content - Retrieve document content

#### Search & Chat Services

- POST /api/v1/search - Perform AI search query
- POST /api/v1/chat/message - Send chat message
- GET /api/v1/chat/sessions - List chat sessions
- DELETE /api/v1/chat/sessions/{id} - Clear chat session
- POST /api/v1/chat/feedback - Submit response feedback

#### Analytics & Monitoring

- GET /api/v1/analytics/overview - Dashboard metrics
- GET /api/v1/analytics/queries - Query statistics
- GET /api/v1/analytics/popular - Popular search terms
- GET /api/v1/health - System health check

### 3.2 API Response Format

```
{  
  "success": true,  
  "data": {},  
  "message": "Operation completed successfully",  
  "timestamp": "2025-09-22T10:30:00Z",  
  "request_id": "uuid-string"  
}
```

### 3.3 Error Handling

```
{  
  "success": false,  
  "error": {  
    "code": "VALIDATION_ERROR",  
    "message": "Invalid query parameters",  
    "details": {}  
  },  
  "timestamp": "2025-09-22T10:30:00Z",  
  "request_id": "uuid-string"  
}
```

---

## 4) Feature Requirements

### 4.1 Administrator Dashboard (React Application)

#### **Content Management Interface**

- **Website Configuration**
  - URL input with validation
  - Crawl depth settings (1-5 levels)
  - Crawl frequency scheduling (hourly, daily, weekly)
  - Robots.txt compliance toggle
  - Sitemap.xml integration

- **Document Management**
  - File upload interface (PDF, DOC, DOCX, TXT, HTML)
  - Document preview and metadata editing
  - Bulk document operations
  - Index status monitoring

## Search Configuration

- **Relevance Tuning**
  - Similarity score thresholds
  - Result ranking parameters
  - Response length controls
  - Citation formatting options
- **Content Filtering**
  - URL pattern inclusion/exclusion
  - Content type filtering
  - Language detection settings
  - Duplicate content handling

## Analytics Dashboard

- **Usage Metrics**
  - Total queries processed
  - Average response times
  - Popular search terms
  - User session statistics
- **Performance Monitoring**
  - System resource usage
  - Crawling job status
  - Error rate tracking
  - Response quality metrics

## Customization Panel

- **Theme Management**
  - Color scheme selection
  - Logo upload and positioning
  - Font family selection
  - Custom CSS injection
- **Branding Controls**
  - Application title and tagline
  - Footer customization
  - White-label options
  - Multi-language UI labels

## 4.2 Embedded Widget (React Component)

### Search Interface

- **localQuery Input**
  - Natural language search box
  - Voice input support (browser API)

- Search suggestions/autocomplete
- Query history (session-based)
- **Results Display**
  - Structured result cards
  - Source attribution with clickable links
  - Relevance score indicators
  - Result snippet highlighting

## Chat Interface

- **Conversation Flow**
  - Message thread display
  - Typing indicators
  - Message timestamps
  - Context preservation across queries
- **Response Formatting**
  - Markdown support for rich text
  - Code syntax highlighting
  - Bullet point and numbered lists
  - Inline citations with hover previews

## Widget Configuration

- **Embedding Options**
    - Multiple integration methods (iframe, script tag, React component)
    - Responsive design breakpoints
    - Z-index and positioning controls
    - Custom event callbacks
  - **Styling System**
    - CSS custom properties
    - Theme inheritance from parent site
    - Mobile-first responsive design
    - Accessibility compliance (WCAG 2.1 AA)
- 

## 5) Technical Requirements

### API Response Times

- Search queries: <10 seconds (95th percentile)
- Chat responses: <10 seconds (95th percentile)
- Document upload: <30 seconds for files up to 50MB
- Dashboard loading: <3 seconds for initial load

### Concurrent Usage

- Support for 50-100 concurrent users
- Queue management for crawling jobs
- Rate limiting: 100 requests per minute per API key
- Session timeout: 30 minutes of inactivity

## **Data Processing**

- **Document Handling**
  - Maximum file size: 50MB per document
  - Supported formats: HTML, PDF, DOC, DOCX, TXT
  - Text extraction with metadata preservation
  - Automatic language detection
- **Content Indexing**
  - Document chunking: 512-1024 token segments
  - Embedding generation: 384-dimensional vectors
  - Incremental indexing with change detection
  - Duplicate content deduplication

## **Storage Requirements**

- **Database Storage**
  - PostgreSQL for structured data
  - Estimated 1GB per 10,000 documents
  - Daily backup capabilities
  - Index optimization for search performance
- **Vector Storage**
  - Chroma vector database
  - Memory-mapped file storage
  - Automatic index rebuilding
  - Similarity search optimization

## **Security & Privacy**

- **API Security**
  - API key-based authentication
  - Request rate limiting
  - Input validation and sanitization
  - SQL injection prevention
- **Data Privacy**
  - Local data processing only
  - No external API calls for LLM inference
  - Configurable data retention policies
  - User session data encryption

---

# **6) User Experience Requirements**

## **6.1 Administrator User Journeys**

### **Initial Setup Flow**

1. **System Configuration:** Set up admin credentials; configure basic system settings; upload branding; test connectivity.

2. **Content Source Setup:** Add website URLs; configure crawling parameters; upload initial documents; verify indexing.
3. **Search Customization:** Adjust relevance parameters; test search; configure response formatting; set up analytics tracking.

### Ongoing Management

1. **Content Maintenance:** Monitor crawl status; review/approve new content; update schedules; manage doc lifecycle.
2. **Performance Optimization:** Review analytics; tune relevance; optimize response times; monitor resources.

## 6.2 End-User Experience

### Search Interaction

1. **Query Entry:** Natural language question input; real-time suggestions; voice input; easy edits.
2. **Result Consumption:** Scan relevant results; verify citations; follow-up queries; share results.

### Chat Conversation

1. **Conversation Initiation:** Clear chat identification; welcome message; example queries; mode switching guidance.
  2. **Ongoing Dialogue:** Natural follow-ups; context retention; source verification; history review.
- 

## 7) Integration Requirements

### 7.1 Embedding Methods

- **JavaScript Widget:** CDN script, config via data attributes, events, minimal DOM footprint
- **React Component:** NPM package, TypeScript types, hooks API, SSR support
- **iframe Integration:** Cross-origin comms, responsive sizing, sandbox

### 7.2 API Integration

- **OpenAPI:** full docs, codegen, interactive testing, SDK generation
  - **Authentication:** API key management, rotation, usage tracking, rate limit config
-

## 8) Deployment & Infrastructure

### 8.1 System Requirements

- **Minimum Hardware:** 4 cores, 8GB RAM (16GB recommended), 100GB SSD, outbound internet for setup
- **OS:** Ubuntu 20.04+ LTS
- **Prereqs:** Docker, Docker Compose, Python 3.9+, Node.js 18+

### 8.2 Installation Process

1. **Pre-deployment:** verify system; network; security baseline; backups
2. **Deployment:** compose up; DB init; model download; SSL; admin creation
3. **Post-deployment:** health checks; initial indexing; performance baseline

#### 8.2.1 Installer & Docker Compose (files)

See

[https://drive.google.com/drive/u/0/folders/1tNawQBWMot5sS0t\\_yzdic0Lg-Ks\\_Kf7m](https://drive.google.com/drive/u/0/folders/1tNawQBWMot5sS0t_yzdic0Lg-Ks_Kf7m)

---

## 9) Success Criteria & Testing

### 9.1 Acceptance Criteria

- **Functional Requirements**
  - All API endpoints respond correctly
  - Search returns relevant results >80% of the time
  - Chat maintains context across conversation
  - Admin dashboard provides complete system control
- **Performance Requirements**
  - API responses under 10-second target
  - System handles 50 concurrent users
  - Crawling processes 10+ pages per minute
  - UI loads and responds within 3 seconds

### 9.2 Quality Assurance

- **Automated Testing:** API integration tests; React unit tests; E2E; perf benchmarking
  - **Manual Testing:** cross-browser; mobile responsiveness; accessibility; security assessment
-

## 10) Development Timeline

### 10.1 Phase 1: Backend API Development

- Core FastAPI application setup
- Database schema and migrations
- Authentication and security layer
- Content crawling and processing APIs
- Search and chat service APIs

### 10.2 Phase 2: Frontend Application Development

- React admin dashboard development
- Embedded widget React components
- API integration and state management
- Responsive design implementation
- Theme and customization system

### 10.3 Phase 3: Integration & Testing

- End-to-end integration testing
  - Performance optimization
  - Security hardening
  - Documentation completion
  - Deployment automation and packaging
- 

## 11) Future Roadmap (v2.0+)

- Multi-tenancy support for SaaS deployment
  - Advanced analytics with custom reporting
  - A/B testing framework for search optimization
  - Advanced security features (RBAC, audit logs)
  - Real-time collaboration features
  - Webhooks; GraphQL API; OAuth/SAML; Third-party integrations
- 

## 12) Personas & User Stories (with Acceptance Criteria)

### 12.1 Personas

- **Admin/Operator:** configures content sources, monitors health/analytics.
- **Integrator (Front-end Dev):** embeds widget into sites/apps.
- **Knowledge Manager:** curates content, approves crawled docs, tunes relevance.

- **End User (Visitor/Agent):** searches or chats to get answers fast.

## 12.2 Top User Stories (MoSCoW)

### Admin/Operator

1. *MUST:* As an Admin, I can add a website to crawl with depth, cadence, and JS rendering toggle so that dynamic sites index correctly.  
**AC:** Given a valid URL, When I save, Then a test render preview and robots/sitemap validation appear; status transitions to **READY**.
2. *MUST:* As an Admin, I can see indexing health, error rates, and last crawl time per source.  
**AC:** Table with status badges; drill-down errors; export CSV.
3. *SHOULD:* As an Admin, I can tune similarity threshold and reranker top-k without redeploy.  
**AC:** Changes persist; new queries reflect settings; versioned config history.

### Integrator

4. *MUST:* As an Integrator, I can embed the widget via <script> or React component with `mode` set to search/chat/auto.  
**AC:** Copy-paste examples; a sandbox page validates API key and connectivity.
5. *SHOULD:* As an Integrator, I can theme the widget to inherit brand tokens (colors, radius, fonts).  
**AC:** CSS variables applied; live preview.

### Knowledge Manager

6. *MUST:* As a KM, I can upload PDFs/DOCs and edit metadata before indexing.  
**AC:** Upload > Preview > Edit > Index pipeline job visible with progress.
7. *SHOULD:* As a KM, I can de-duplicate near-identical content.  
**AC:** Suggested merges with similarity score; one-click dedupe.

### End User


8. *MUST:* As a User, I can ask a question and get an answer with citations to original sources.  
**AC:** Response shows at least 3 citations when available; clicking opens source.
  9. *MUST:* As a User, I can give thumbs-up/down with reason tags.  
**AC:** Feedback recorded; Admin analytics reflect aggregates.
  10. *SHOULD:* As a User, I can switch between Search results and Chat answer in the same panel.  
**AC:** Tabbed UI or toggle; shared query state.
-

### 13) Risks & Mitigations


#	Risk	Likelihood	Impact	Mitigation	Owner	Trigger/Monitor
R1	Dynamic sites not captured by static crawler	Medium	High	Headless rendering via Playwright; preview endpoint; per domain render budget	Backend TL	Crawl error rate >5% for JS pages
R2	Long LLM latencies on CPU-only hosts	Medium	High	Default to Mistral 7B, response streaming, caching; document HW sizing	Platform Lead	p95 latency >10s
R3	Vector DB bloat & slow recall	Medium	Medium	Chunk size per type, pruning policy, HNSW tuning, nightly compaction	Data Eng	Recall time >500ms or DB size growth >10%/wk
R4	Hallucinations on sparse corpora	Medium	High	Strict citation requirement, low temperature, reranker, claim-check heuristics	QA + NLP	Feedback downvotes >15%
R5	Compliance/privacy concerns	Low	High	Local inference only, configurable retention, PHI/PII filters, audit logs v2	Sec/Compliance	Periodic audits
R6	Rate-limit abuse on public widgets	Medium	Medium	Per-key quotas + IP throttling, CAPTCHA on burst, CDN caching	DevOps	429 spikes/DoS alerts
R7	Model/embedding version drift	Medium	Medium	Versioned index; dual-write and flip read alias; blue/green vectors	Data Eng	Model version change PR
R8	Single-node outages	Low	High	Compose profiles for HA, health checks, auto backups, documented restore	DevOps	Health probe failures
R9	Content license/copyright	Low	Medium	Respect robots, honor crawl allowlist, manual approval queue	KM Lead	Legal review flags

## 14) Sequence Diagrams

### 14.1 Search/Chat (unified)




### 14.2 Crawl/Index



# 15) High-Level Data Model

## 15.1 ER Diagram



## 15.2 Key Fields (illustrative)

ORGANIZATION(id, name, plan, created\_at)

USER(id, org\_id, email, role[admin|km|viewer], password\_hash, created\_at)

APIKEY(id, org\_id, key\_hash, label, rate\_limit\_pm, created\_at)

CRAWLTARGET(id, org\_id, base\_url, depth, cadence, headless[on|off|auto], allowlist, denylist, status)

CRAWLJOB(id, target\_id, started\_at, finished\_at, pages\_fetched, errors, status)

DOCUMENT(id, org\_id, source\_url, title, lang, checksum, status, last\_indexed\_at)

CHUNK(id, doc\_id, ord, text, token\_count, hash)

EMBEDDING(id, chunk\_id, model, dim, vector\_ref)

CHATSESSION(id, org\_id, user\_id|null, mode[search|chat|auto], started\_at, meta)

MESSAGE(id, session\_id, role[user|assistant|system], content, citations[], latency\_ms, created\_at)

QUERYLOG(id, org\_id, apikey\_id, query, mode, p95\_latency, result\_count, timestamp)

FEEDBACK(id, message\_id, vote[up|down], reasons[], note, created\_at)

ANALYTICSDAY(date, org\_id, queries, avg\_latency\_ms, thumbs\_up\_rate, errors)

SETTING(id, org\_id, key, value\_json, version, updated\_by)

BRANDASSET(id, org\_id, logo\_url, colors\_json, css\_overrides)

---

## 16) Defaults & Tunables (v1)

- **Embeddings:** [sentence-transformers/all-MiniLM-L6-v2](#) (384-d)
- **Chunking:** HTML 800, PDF/DOC 600, MD/TXT 900; overlap 80
- **Retriever:** Chroma cosine; [topK=24](#)
- **Re-ranker:** BM25 default; optional mono-\* (off by default)
- **Context Window:** top 10 chunks post-rerank (~2,500 tokens)
- **LLM:** [mistral:7b](#) default; [llama2:7b](#) alt; temperature 0.2; streaming on

- **Citations:** require  $\geq 3$  unique docs if available; fail closed if none
  - **Dedup:** LSH on chunk hash; threshold 0.92
  - **Cache:** Redis hash(query+config); TTL 10 min; stream-friendly
  - **Headless Rendering:** Auto when low DOM text or missing selectors; 3s/page budget
  - **Crawl Cadence:** Daily default; hourly/weekly per target; backoff on fails
  - **Rate Limits:** 100 rpm/key; burst 20; 429 w/ **Retry-After**
- 

## 17) Release, Canary & Rollback

- **Blue/Green** deploy; Nginx upstream points to **active** label
  - **Canary:** route 10% via header **X-Canary: 1** or IP hash
  - **Promote** when error  $<1\%$  and p95  $<10\text{s}$  for 30 min
  - **Rollback:** flip upstream; Alembic down for DB; vector dual-index alias
  - **Feature Flags:** **RERANKER\_ON**, **STRICT\_CITATIONS**, **HEADLESS\_AUTO**
- 

## 18) Test Plan Additions

- **RAG Correctness:** 100 Q/A golden set; MRR/NDCG; cite coverage
  - **Latency:** k6 50 VU / 10 min; assert p95  $<10\text{s}$
  - **Crawler:** JS-heavy sample; preview vs indexed text similarity  $>0.8$
  - **Security:** OWASP Top-10; fuzz inputs
  - **Resilience:** kill/restart containers; recover  $<30\text{s}$
- 

## 19) Backlog: GitHub/Jira (Owners & Timeboxes)

Roles: **BE** (Backend Lead), **FE**, **DE** (Data Eng), **DEV**, **QA**, **DO** (DevOps).  
Two-week sprints  $\times 8$ .

### Epics

- **E1** Platform Foundations (Infra, API skeleton, Auth)
- **E2** Crawl & Ingestion
- **E3** RAG/Search/Chat Core
- **E4** Admin App
- **E5** Embed Widget
- **E6** Analytics & Feedback
- **E7** Packaging, CI/CD, Installer
- **E8** QA, Perf, Security

### Issues (sample)

Key	Title	Epic	Owner	Est.	Depends
P-1	FastAPI skeleton + auth + OpenAPI	E1	BE		—
P-2	DB schema + Alembic migrations	E1	BE		P-1
P-3	Redis cache scaffolding	E1	DEV		P-1
C-1	Crawler service (Scrapy+Playwright)	E2	DEV		P-2
C-2	Robots/sitemap resolver	E2	DEV		C-1
C-3	Crawl preview endpoint	E2	BE		C-1
I-1	Ingestion pipeline (clean/split)	E2	DE		P-2
I-2	Embedding generation + Chroma client	E2	DE		I-1
Q-1	Retriever + topK search	E3	DE		I-2
Q-2	Reranker (BM25/mono-*)	E3	DE		Q-1
Q-3	Prompt builder + citation rules	E3	BE		Q-2
Q-4	Chat endpoint (streaming)	E3	BE		Q-3
A-1	Admin shell (Vite+React+TS)	E4	FE		P-1
A-2	Content sources UI + preview	E4	DEV		C-3
A-3	Relevance tuning UI	E4	DEV		Q-2
A-4	Analytics dashboard	E4	DEV		E6
W-1	Widget (search/chat/auto)	E5	FE		Q-4
W-2	Theming (CSS vars) + docs	E5	DEV		W-1
X-1	Feedback API + UI	E6	BE		Q-4
X-2	Analytics ETL + PG views	E6	DE		X-1
PK-1	Docker Compose + Makefile	E7	BE		P-1
PK-2	Installer script + Nginx	E7	BE		PK-1

Key	Title	Epic	Owner	Est.	Depends
PK-3	Blue/green + canary doc	E7	BE		PK-2
QA-1	E2E tests (Playwright/Cypress)	E8	QA		E4,E5
QA-2	Load tests (k6)	E8	QA		Q-4
SEC-1	Security checklist	E8	QA		P-1

**Definition of Done:** code + tests + docs; CR passed; merged to main; CI green; feature behind flag if risky.

**Sprint Cadence:** 2-week sprints × 8; demo at Sprints 2/4/6; RC at Sprint 8.

---

## 20) Low-Fi Mockups / Prototype

### 20.1 Admin — Dashboard (ASCII)

```
+-----+
| Sidebar | Analytics | Jobs | Errors |
+-----+-----+
| ▶ Home | [ Cards: Queries Today | p95 ] | |
| ▶ Crawl | Thumbs↑ rate | Fail % |
| ▶ Docs | [ Chart: Queries over time ] |
| ▶ RAG | [ Table: Top sources | 5xx | 4xx ] |
| ▶ Theme |
+-----+
```

### 20.2 Admin — Crawl Source Detail

URL: [ https://example.com ] Depth: [2] Cadence: [Daily]

Headless: [Auto] Robots: [✓] Sitemap: [✓]

[ Preview Render ] [ Test Crawl ] [ Save ]

Allowlist: \*.example.com/docs/\*

Denylist: \*.example.com/private/\*

Status: READY Last Crawl: 2025-09-22 22:10 Errors: 0

## 20.3 Widget — Unified Search/Chat

[ Ask anything about ACME... (  ) ] (Search | Chat | Auto)

---

Answer: "Here's a summary..." [  ] [  ]

Sources: [Doc A] [Doc B] [Doc C]

[ Show results tab ] [ Follow-up question... ]

## 20.4 Embed Examples

<!-- Script tag (Auto mode) -->

```
<script src="/widget/embed.js"></script>
```

```
<script>
```

```
AISeachChat.init({
```

```
    apiBase: "https://your-domain.example/api/v1",
```

```
    apiKey: "public-embeddable-key",
```

```
    mode: "auto",
```

```
    container: "#ai-widget",
```

```
    theme: { primary: "#0052cc", radius: 10 }
```

```
});
```

```
</script>
```

```
<div id="ai-widget"></div>
```

```
// React (TypeScript)
```

```
<AISeachChatWidget
```

```
apiBase="/api/v1"  
  
apiKey={env.PUBLIC_WIDGET_KEY}  
  
mode="search"  
  
onFeedback={(res, vote) => console.log({res, vote})}  
  
/>
```

---

## 21) Admin & API Quickstart (Docs excerpt)

### 21.1 Admin

1. Install via [install.sh](#).
2. Login [/admin](#) with default admin (change password).
3. Add **Crawl Target**; run preview; save & start crawl.
4. Upload a test PDF; verify documents appear.
5. Tune **RAG** thresholds; test queries in **Playground**.

### 21.2 API

POST /api/v1/search

```
{ "query": "Refund policy for EU?", "top_k": 10 }
```

Response includes: [answer](#), [citations\[\]](#), [latency\\_ms](#), [request\\_id](#).

---

## 22) Open Questions

- What do we think about integrate MCP? (*Maybe in v2*)
  - Default LLM for CPU-only: **mistral:7b?** (*Proposed: yes*)
  - Tenancy: single-org in v1, multi-tenant in v2? (*Proposed: single-org v1*)
- 

## 23) Document Approval

- Technical Architecture Review
- Product Management Approval
- Development Team Sign-off

- Quality Assurance Review
- Stakeholder Final Approval